

HW4: Word Segmentation

Virgile Audi
vaudi@g.harvard.edu
Nicolas Drizard
nicolasdrizard@g.harvard.edu
Github: virgodi/cs287/HW5

April 17, 2016

1 Introduction

The goal of this assignment is to tackle the NLP task of identifying and labeling contiguous segments of text. We will use sequence models and a dynamic programming method to find the best scoring sequence.

2 Problem Description

The idea is here to label continuous sequence of words with BIO tagging of different entities. The entities are the following:

1. PER: a person
2. LOC: a location
3. ORG: an organization
4. MISC:

Furthermore, this tagging method identifies the continuous group of words belonging to the same entity: the prefix B stop the current tag and begins a new one whereas the prefix I continues adding to the previous tag. However, in our solution we just cared about predicting the entity tag and then we were grouping the contiguous predictions into the same entity because the training text does not contain any B-tag.

3 Model and Algorithms

We used three different methods to solve this problem. The first two are the equivalent of first the Naive Bayes and second the logistic regression from text classification tasks. The last one introduces a customized way to train a neural architecture for this task.

3.1 Hidden Markov Model

We implement here a standard first order hidden Markov Model. The hidden states are the tags and the observed states are the features we built (word counts, capitalization...). The model can be represented with the following graphical model and requires two distributions: emission and transition.

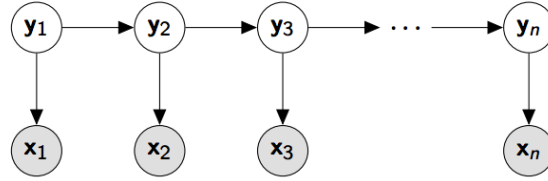


Figure 1: Graphical model of 1st order HMM with one feature

We represent the two distributions with multinomial as they model feature counts. As a result, we can infer them simply with the maximum likelihood estimator:

$$p(x_i = \delta(f) | y_i = \delta(c)) = \frac{F_{f,c}}{F_{.,c}}$$

$$p(y_i = \delta(c_i) | y_{i-1} = \delta(c_{i-1})) = \frac{T_{c_{i-1},c_i}}{T_{c_{i-1},.}}$$

with T_{c_{i-1},c_i} the counts of class c_{i-1} preceding class c_i and $F_{f,c}$ the counts of emission f with class c .

If we consider multiple features, then we still assume that the features are independent with each other (it's the main assumption in the Naive Bayes approach also). Only the emission distribution is changed and we can combine the probability together:

$$p(x_i = (\delta(f_1), \delta(f_2)) | y_i = \delta(c)) = p(x_i = \delta(f_1) | y_i = \delta(c)) p(x_i = \delta(f_2) | y_i = \delta(c)) = \frac{F_{f_1,c}}{F_{.,c}} \frac{F_{f_2,c}}{F_{.,c}}$$

3.2 Maximum-Entropy Markov Model

Next, we implemented a Maximum-Entropy Markov Model. The objective of the MEMM is to evaluate at each time step a distribution over the possible tags using features of the current word, denoted as $feat(x_i)$ and the tag of the previous word, c_{i-1} , using multi-class logistic regression, i.e.

$$p(y_i | y_{i-1}, feat(x_i)) = \text{softmax}([feat(x_i), c_{i-1}] \mathbf{W} + \mathbf{b})$$

3.3 Viterbi algorithm

The search algorithm that we implemented is the dynamic programming algorithm named after Andrew Viterbi. Its main difference with a greedy approach is that it evaluates at every step and for every previous state, the best possible next step. This guarantees a solution closer to the true optimal solution. The pseudo-code of the algorithm is given by:

```
procedure VITERBIWITHBP
   $\pi \in \mathbb{R}^{n+1 \times \mathcal{C}}$  initialized to  $-\infty$ 
   $bp \in \mathcal{C}^{n \times \mathcal{C}}$  initialized to  $\epsilon$ 
   $\pi[0, \langle s \rangle] = 0$ 
  for  $i = 1$  to  $n$  do
    for  $c_{i-1} \in \mathcal{C}$  do
      compute  $\hat{y}(c_{i-1})$ 
      for  $c_i \in \mathcal{C}$  do
         $score = \pi[i-1, c_{i-1}] + \log \hat{y}(c_{i-1})_{c_i}$ 
        if  $score > \pi[i, c_i]$  then
           $\pi[i, c_i] = score$ 
           $bp[i, c_i] = c_{i-1}$ 
  return sequence from  $bp$ 
```

3.4 Structured Perceptron

The final model, we implemented is the structure perceptron train algorithm. The way the model is trained uses the Viterbi search algorithm, presented above. At each epoch, we uses Viterbi to predict the highest scored sequence given the state of the model. We can then find the timesteps where the actual sequence for the given sentence and the predicted one differ and compute at each of these time steps, the gradient of a hinge type loss. These gradients have a -1 entry on the true class for this given word, and a 1 on the predicted class by the model. We can then propagate these gradients in the network, and update the weights with a learning rate that can be tuned.

The model itself is similar to the model of the MEMM without the final logsoftmax layer.

4 Experiments

4.1 Feature Engineering

The original paper suggests several features to use. We focus on the word counts and a capitalization feature. We defined our capitalization feature as follow:

1. 1 : word in low caps;
2. 2 : whole word in caps;
3. 3 : first letter in cap;
4. 4 : one cap in the word;

5. 5 : other

We then produced an embedding of the word counts using a pre-trained version.

We also used the Python "pattern.en" package to extract Part-of-Speech (PoS) features. The packages generates 41 features to which we added special feature for the opening and closing tabs `<s>` and `< \s>`.

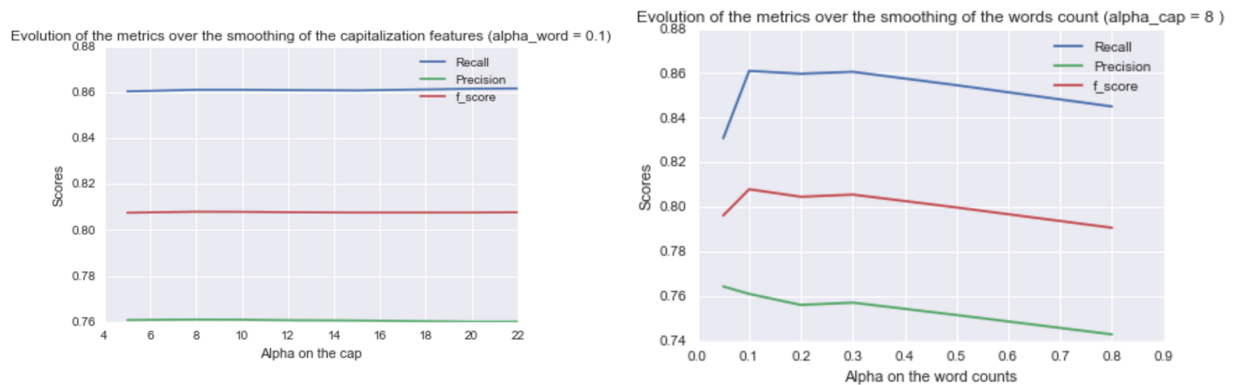
4.2 Model Evaluation

As used in the Kaggle competition, we used the f-score with the precision and recall measure to evaluate our model while tuning the hyperparameters. A positive prediction stands for a label (in the notation of the task, everything which is not the **O** tag):

1. recall: ratio of the true positive predictions among the positives tags in the correct sequence
2. precision: ratio of the true positive predictions among the positive predictions,
3. f-score (with $\beta = 1$): harmonic mean of the precision and the recall, i.e. $f_1 = \frac{2pr}{p+r}$

4.3 Hidden Markov Model

There is only the smoothing parameter α and eventually feature selection here to tune here. We evaluate the impact of adding more features and run experiments with different alpha values to tune them. One important details is to make sure to use a specific smoothing parameter for each distribution, i.e a smoothing parameter may be applied to the transition matrix but also to the emission matrix of each different feature. Each of this distribution has a different tail and need a different smoothing. For instance, the transition matrix need a very small α (around 0.1) because we are pretty confident in it but the capitalizations feature need one much bigger (around 20) because the counts are already high.



We notice that the model is less sensitive to the changes of the smoothing parameter on the capitalization feature as on the word counts. This is pretty reasonable as the feature counts are much higher in the capitalization feature than in the word counts. Tuning this parameter provides

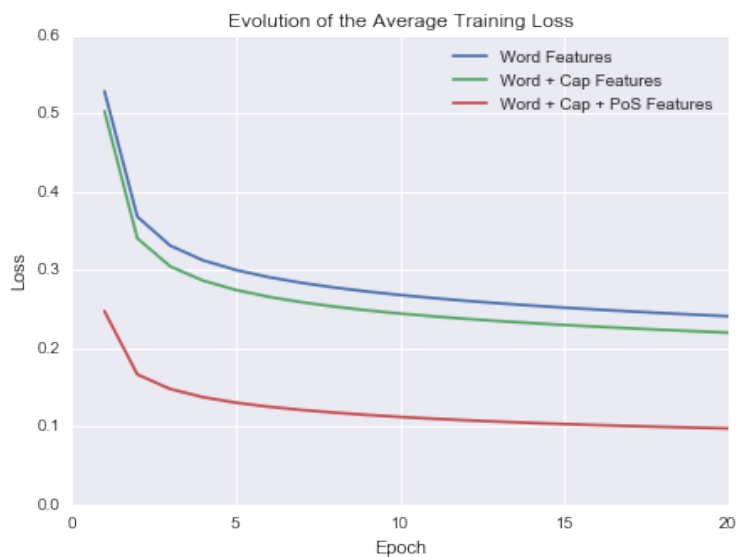
a model with a f-score of **0.808**. Using only the word counts features provide a best f-score of **0.764**.

Adding the part-of-speech tagging feature increased the performance of the model. We got on the dev set a f-score of **0.843** We obtained a Kaggle score on the test set with the three different features of :

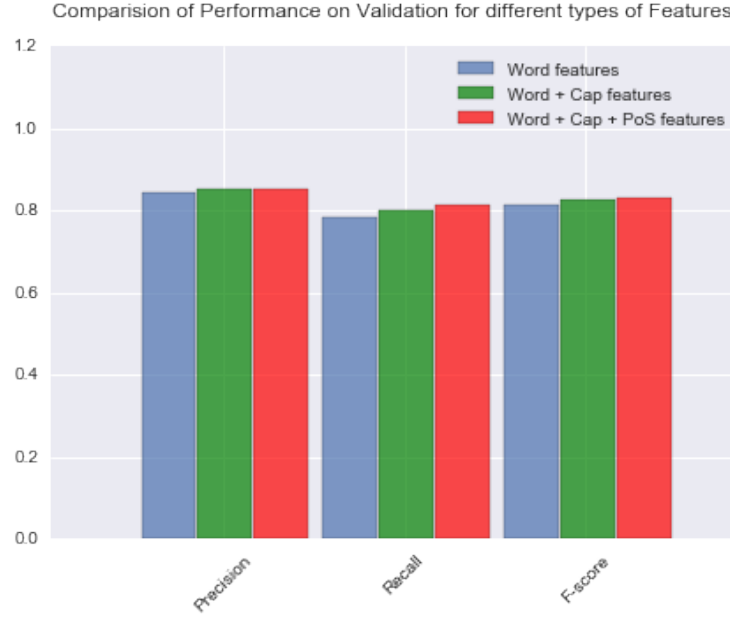
$$K_{HMM} = 0.54392$$

4.4 Maximum-Entropy Markov Model

We coded the MEMM using the nn module and trained using stochastic gradient descent. We also used the Glove embeddggins using a lookup table. As for the HMM, we used two different sets of features, i.e. the words and the words and capitalisation of the words. We also added the Part of Speech features that were evaluated using the python package "pattern.en" in order to gain some time. We observed that the training algorithm converges quite rapidly, and that if adding caps to the features helped decrease the loss, the impact was not as strong as expected. On the other hand, adding PoS features impacted greatly the loss. Nevertheless, we trained the model on 20 epochs in order to learn the embeddings for the <s> and <\s> "words" added during pre-processing.



We evaluated the performance of these two models using the f-score presented above:



Adding extra features yielded better results on both Precision and Recall and therefore on the f-score. But as we expected from the small differences in loss, we did not observe an important increase on the f-score using cap features. We were nevertheless surprised to see that the impact on loss using PoS features did not translate on the f-score. These results were later confirmed on the test set, as the kaggle score obtained for these two models were:

$$K_{nocaps} = 0.52057 \quad \text{and} \quad K_{caps} = 0.55482 \quad K_{PoS} = 0.57121$$

which are both slightly better than the results of the HMM.

4.5 Structured Perceptron

We implemented the structured perceptron with the idea described in the model: weighting up the true edges in the lattice and down the incorrectly predicted. However we did not observe convincing results on our model, especially the f-score on the dev set was not increasing over the epochs but simply oscillating randomly around 0.72, which is not so bad but still less than what we obtained from the two other model.

We first tried a simple training version (our first train function in sp.lua) where for each timestep with a wrong prediction we do one forward/backward. The input is the right tag and we use a gradient with -1 on the right tag and 1 on the wrong one. We also coded an advanced version which was treating the two wrong edges of the lattice for each error (in our second train function) but did not observe the expected result.

5 Conclusion

We were disappointed to not be able to get the performance of the structured perceptron to the levels of the hidden and maximum-entropy markov models on the task of finding and labeling

named-entities in text. Due to our difficulties at implementing the perceptron, we did not get the chance of implementing the NNMEM or add more features. Nevertheless, looking at the impact of the few extra features implemented and the performance of the multi-class logistic regression, we believe that future work in that direction would yield better results.

Appendices

Preprocessing:

```
1 import numpy as np
2 import h5py
3 import re
4 import pattern.en
5 import sys
6 import argparse
7
8 from itertools import product
9
10
11 def get_tag2index():
12     # Tags mapping
13     tag2index = {}
14
15     with open('data/tags.txt', 'r') as f:
16         for line in f:
17             line_split = line[:-1].split(' ')
18             tag2index[line_split[0]] = int(line_split[1])
19
20     # Adding tags for end/start of sentence
21     tag2index['<t>'] = 8
22     tag2index['<\t>'] = 9
23     return tag2index
24
25
26 def get_pos2index():
27     '''
28     Part of speech tagging tags to feature index mapping
29     '''
30     # mapping for the POS tags
31     tags = ['CC', 'CD', 'DT', 'EX', 'FW', 'IN', 'JJ', 'JJR', 'JJS', 'LS',
32            'MD',
33            'NN', 'NNS', 'NNP', 'NNPS', 'PDT', 'POS',
34            'PRP', 'PRP$', 'RB', 'RBR', 'RBS', 'RP', 'SYM', 'TO', 'UH',
35            'VB',
```

```

34         'VBZ', 'VBP', 'VBD', 'VBN', 'VBG', 'WDT', 'WP', 'WP$', 'WRB
35         ,
36         '.', ',', ':', '(', ')']
37     pos2index = {k: v+1 for v, k in enumerate(tags)}
38     return pos2index
39
40
41 def count_elements(filename, tags=True):
42     # Counting the number of elements to stored (ie num_words +
43     # 2*num_sentences)
44     num_words = 0
45     num_sentences = 0
46     with open(filename, 'r') as f:
47         for line in f:
48             if tags:
49                 line_split = line[:-1].split('\t')
50             else:
51                 line_split = line[:-1].split(' ')
52             # Case blank
53             if len(line_split) == 1:
54                 num_sentences += 1
55             else:
56                 num_words += 1
57
58     return num_words, num_sentences
59
60
61 def get_cap_feature(word):
62     # Return the caps feature for the given word
63     # 1 – low caps; 2 – all caps; 3 – first cap; 4 – one cap; 5 – other
64     if len(word) == 0 or word.islower() or re.search('[.?\"', word)
65     :
66         feature = 1
67     elif word.isupper():
68         feature = 2
69     elif len(word) and word[0].isupper():
70         feature = 3
71     elif sum([w.isupper() for w in word]):
72         feature = 4
73     else:
74         feature = 5
75     return feature
76

```



```

77 def get_tokenized_sentences(filename, tags=True):
78     # Build the part of speech tags
79     with open(filename, 'r') as f:
80         text = []
81         for line in f:
82             if tags:
83                 line_split = line[:-1].split('\t')
84             else:
85                 line_split = line[:-1].split(' ')
86             if len(line_split) != 1:
87                 text.append(line_split[2])
88
89     return pattern.en.tag(' '.join(text))
90
91
92 def build_input_matrix(filename, num_rows, tag2index, pos2index, tags=
True, word2index=None, memm = False):
93     # Building input matrix with columns: (id, id_in_sentence, id_word,
id_caps, id_token, id_tag)
94     # caps feature:
95     # 1 – low caps; 2 – all caps; 3 – first cap; 4 – one cap; 5 – other
96     # Tags: if correct solution given (ie 4th column)
97     # word2index: if use of previously built word2index mapping
98
99     # Features for starting/ending of sentence (3 last columns)
100    # For the POS tag, we use the same as a point (index 36)
101    # initialization
102    input_matrix = np.zeros((num_rows, 6), dtype=int)
103    if memm == False:
104        input_matrix[0] = [1, 1, 1, 1, 36, 8]
105        start = [1, 1, 36, 8]
106        end = [2, 1, 36, 9]
107    else:
108        input_matrix[0] = [1,1,word2index['<s>'],1,36,8]
109        start = [word2index['<s>'],1,36, 8]
110        end = [word2index['<\s>'],1,36, 9]
111    row = 1
112
113    # Get the POS token
114    tokenized_sentences = get_tokenized_sentences(filename, tags=tags)
115    pos_i = 0
116
117    # Boolean to indicate if a sentence is starting
118    starting = False

```

```

119     # Boolean if a mapping is defined (last element of the mapping is
120     # for
121     # unknown words)
122     if word2index == None:
123         test = False
124         word2index = {'<s>': 1, '<\s>': 2}
125         id_word = 3
126     else:
127         test = True
128     with open(filename, 'r') as f:
129         for line in f:
130             if tags:
131                 line_split = line[:-1].split('\t')
132             else:
133                 line_split = line[:-1].split(' ')
134             if starting == True:
135                 # Start of sentence
136                 input_matrix[row, 0] = input_matrix[row-1, 0] + 1
137                 input_matrix[row, 1] = 1
138                 input_matrix[row, 2:] = start
139                 row += 1
140                 starting = False
141             if len(line_split) == 1:
142                 # End of sentence
143                 input_matrix[row, :2] = input_matrix[row-1, :2] + 1
144                 input_matrix[row, 2:] = end
145                 row += 1
146                 starting = True
147             else:
148                 # Indexing
149                 input_matrix[row, 0] = input_matrix[row-1, 0] + 1
150                 input_matrix[row, 1] = int(line_split[1]) + 1
151                 # Build cap feature
152                 word = line_split[2]
153                 input_matrix[row, 3] = get_cap_feature(word)
154                 # Build pos feature
155                 pos_tag = tokenized_sentences[pos_i][1].split('-')[0]
156                 if pos_tag in pos2index.keys():
157                     input_matrix[row, 4] = pos2index[pos_tag]
158                 else:
159                     input_matrix[row, 4] = len(pos2index) + 1
160                 pos_i += 1
161             # Build word count feature
162             word_clean = word.lower()

```

```

163         if not test:
164             if word_clean not in word2index:
165                 word2index[word_clean] = id_word
166                 id_word += 1
167             input_matrix[row, 2] = word2index[word_clean]
168         else:
169             # Unseen word during train
170             if word_clean not in word2index:
171                 input_matrix[row, 2] = len(word2index)
172             else:
173                 input_matrix[row, 2] = word2index[word_clean]
174         if tags:
175             input_matrix[row, 5] = tag2index[line_split[3]]
176         row += 1
177     # Add special word if training
178     if not test:
179         word2index['<unk>'] = len(word2index)+1
180     if tags:
181         return input_matrix, word2index
182     else:
183         return input_matrix[:, :5], word2index
184
185 #Function that formats the output of the previous function in order to
    run MEMM:
186 def input_mm_pos(matrix):
187
188     nwords = matrix.shape[0]
189
190     res = np.zeros((nwords,1 + 9 + 5 + 43 + 1),dtype = int)
191
192     res[:,0] = matrix[:,2]
193
194     for i in range(nwords):
195         tag_1_hot = np.zeros(9)
196         tag_1_hot[matrix[i,5]-1] = 1
197         tag_1_hot_cap = np.zeros(5)
198         tag_1_hot_cap[matrix[i,3]-1] = 1
199         tag_1_hot_pos = np.zeros(43)
200         tag_1_hot_pos[matrix[i,4]] = 1
201         res[i,1:10] = tag_1_hot
202         res[i,10:15] = tag_1_hot_cap
203         res[i,15:58] = tag_1_hot_pos
204     res[:,58] = matrix[:,5]
205     return res
206

```

```

207
208 def train_hmm(input_matrix, num_features, num_pos, num_tags):
209     # Emission word_count matrix:
210     # size (num_words, num_tags)
211     # row: observation / colum: tag
212     # (un-normalized if smoothing required)
213     emission_w = np.zeros((num_features, num_tags), dtype=int)
214
215     # Emission caos_count matrix:
216     # size (5, num_tags)
217     # row: observation / colum: caps
218     # (un-normalized if smoothing required)
219     emission_c = np.zeros((5, num_tags), dtype=int)
220
221     # Emission pos_count matrix:
222     # size (5, num_tags)
223     # row: observation / colum: pos tag
224     # (un-normalized if smoothing required)
225     emission_p = np.zeros((num_pos, num_tags), dtype=int)
226
227     # Building
228     for r in input_matrix:
229         emission_w[r[2]-1, r[5]-1] += 1
230         emission_c[r[3]-1, r[5]-1] += 1
231         emission_p[r[4]-1, r[5]-1] += 1
232
233     # Transition matrix
234     # size (num_tags, num_tags)
235     # row: to / colum: from
236     # (un-normalized if smoothing required)
237     transition = np.zeros((num_tags, num_tags), dtype=int)
238     for i in xrange(input_matrix.shape[0] - 1):
239         transition[input_matrix[i+1, 5]-1, input_matrix[i, 5]-1] += 1
240
241     return emission_w, emission_c, emission_p, transition
242
243
244 def main(arguments):
245     # Args
246     global args
247     parser = argparse.ArgumentParser(
248         description=__doc__,
249         formatter_class=argparse.RawDescriptionHelpFormatter)
250
251     parser.add_argument('-f', default='data/features.hdf5',

```

```

252                                     type=str , help='Filename to save data')
253 args = parser.parse_args(arguments)
254 filename = args.f
255
256 # Train
257 pos2index = get_pos2index()
258 tag2index = get_tag2index()
259 num_words, num_sentences = count_elements('data/train.num.txt')
260 num_rows = num_words + 2*num_sentences
261 input_matrix_train, word2index = build_input_matrix('data/train.num
    .txt',
262                                                     num_rows,
263                                                     tag2index,
264                                                     pos2index)
265
266 # Building the count matrix
267 num_tags = len(tag2index)
268 num_features = len(word2index)
269 num_pos = len(pos2index) + 1
270 emission_w, emission_c, emission_p, transition = train_hmm(
    input_matrix_train,
271                                                     num_features
272                                                     ,
273                                                     num_pos
274                                                     ,
275                                                     num_tags
276                                                     )
277
278 # Dev & test
279 num_words, num_sentences = count_elements('data/dev.num.txt')
280 # Miss 1 blank line at the end of the file for the dev set
281 num_rows = num_words + 2*num_sentences + 1
282 input_matrix_dev, word2index = build_input_matrix('data/dev.num.txt
    ',
283                                                     num_rows,
284                                                     tag2index,
285                                                     pos2index,
286                                                     word2index=
287                                                     word2index)
288
289 num_words, num_sentences = count_elements('data/test.num.txt',
290                                                     tags=False)
291
292 num_rows = num_words + 2*num_sentences
293 input_matrix_test, word2index = build_input_matrix('data/test.num.
    txt',

```

```

286                                     num_rows,
287                                     tag2index,
288                                     pos2index,
289                                     tags=False,
                                     word2index=
                                     word2index)

290
291     # Saving pre-processing
292     with h5py.File(filename, "w") as f:
293         # Model
294         f['emission_w'] = emission_w
295         f['emission_c'] = emission_c
296         f['emission_p'] = emission_p
297         f['transition'] = transition
298
299         f['input_matrix_train'] = input_matrix_train
300         f['input_matrix_dev'] = input_matrix_dev
301         f['input_matrix_test'] = input_matrix_test
302
303
304     if __name__ == '__main__':
305         sys.exit(main(sys.argv[1:]))

```

Hidden Markov Model:

```

1  — Documentation:
2  — ——— How to call it from the command line?
3  — For example:
4  — $ th count_based.lua -N 5
5  — Other argument possible (see below)
6  —
7  — ——— Is there an Output?
8  — By default, the predictions on the test set are saved in hdf5 format
   as classifier .. opt.f
9
10 — Only requirements allowed
11 require("hdf5")
12 require 'helper.lua';
13
14 cmd = torch.CmdLine()
15
16 — Cmd Args
17 cmd:option('-datafile', 'data/words_feature.hdf5',
18            'Datafile with features in hdf5 format')
19 cmd:option('-alpha_t', 0.1, 'Smoothing parameter alpha in the
   transition counts')

```

```

20 cmd:option('--alpha_w', 0.1, 'Smoothing parameter alpha in the word
    counts')
21 cmd:option('--alpha_c', 8, 'Smoothing parameter alpha in the caps counts
    ')
22 cmd:option('--alpha_p', 2, 'Smoothing parameter alpha in the pos counts
    ')
23 cmd:option('--test', 0, 'Boolean (as int) to ask for a prediction on
    test, will be saved in submission in hdf5 format')
24 cmd:option('--datafile_test', 'submission/v_seq_hmm', 'Smoothing
    parameter alpha in the word counts')
25 cmd:option('--nfeatures', 2, 'Number of type of features to use')
26 cmd:option('--cv', 0, 'Boolean (as int) to run a cross validation
    pipeline')
27
28
29
30 — Formating as log-probability and smoothing the input
31 function format_matrix(matrix, alpha)
32     local formatted_matrix = matrix:clone():type('torch.DoubleTensor')
33     formatted_matrix:add(alpha)
34     — Normalize
35     local norm_mat = torch.expandAs(formatted_matrix:sum(1),
        formatted_matrix)
36     formatted_matrix:cdiv(norm_mat)
37     return formatted_matrix:log()
38 end
39
40 — log-scores of transition and emission
41 — corresponds to the vector y in the lecture notes
42 — i: timestep for the computed score
43 function score_hmm(observations, i, emissions, transition, C, nfeatures
    )
44     local observation_emission = torch.zeros(C)
45     for k=1,nfeatures do
46         — print(i,k)
47         — print(emissions[k][observations[{i,k}]]))
48         observation_emission:add(emissions[k][observations[{i,k}]]))
49     end
50     observation_emission = observation_emission:view(C, 1):expand(C, C)
51     — NOTE: allocates a new Tensor
52     return observation_emission + transition
53 end
54
55 — Viterbi algorithm.
56 — observations: a sequence of observations, represented as integers

```

```

57 — logscore: the edge scoring function over classes and observations in
    a history-based model
58 function viterbi(observations, logscore, emissions, transition,
    nfeatures)
59     local y
60     — Formating tensors
61     local initial = torch.zeros(transition:size(2), 1)
62     — initial started with a start of sentence: <t>
63     initial[{8,1}] = 1
64     initial:log()
65
66     — number of classes
67     C = initial:size(1)
68     local n = observations:size(1)
69     local max_table = torch.Tensor(n, C)
70     local backpointer_table = torch.Tensor(n, C)
71
72     — first timestep
73     — the initial most likely paths are the initial state distribution
74     — NOTE: another unnecessary Tensor allocation here
75     local init_pred = initial:clone()
76     for i=1,nfeatures do
77         init_pred:add(emissions[i][observations[{1,i}]]))
78     end
79     local maxes, backpointers = init_pred:max(2)
80     max_table[1] = maxes
81
82     — remaining timesteps ("forwarding" the maxes)
83     for i=2,n do
84         — precompute edge scores
85         y = logscore(observations, i, emissions, transition, C,
            nfeatures)
86         scores = y + maxes:view(1, C):expand(C, C)
87
88         — compute new maxes (NOTE: another unnecessary Tensor
            allocation here)
89         maxes, backpointers = scores:max(2)
90
91         — record
92         max_table[i] = maxes
93         backpointer_table[i] = backpointers
94     end
95     — follow backpointers to recover max path
96     local classes = torch.Tensor(n)
97     maxes, classes[n] = maxes:max(1)

```



```

98     for i=n,2,-1 do
99         classes[i-1] = backpointer_table[{i, classes[i]}]
100     end
101
102     return classes
103 end
104
105 — Prediction pipeline
106 function predict(observations, emissions, transition, alphas, nfeatures
107 )
108     — Formating model parameters (log and alpha smoothing)
109     — Alphas is a tensor : {alpha_t, alpha_w, alpha_c}
110     emissions_cleaned = {}
111     for i=1,nfeatures do
112         emissions_cleaned[i] = format_matrix(emissions[i], alphas[i+1])
113     end
114     local transition_cleaned = format_matrix(transition, alphas[1])
115     return viterbi(observations, score_hmm, emissions_cleaned,
116         transition_cleaned, nfeatures)
117 end
118 — Cross validation pipeline
119 function cross_validation(observations, emissions, transitions,
120     true_classes,
121     alphas_table, alpha_t)
122     — alphas_table is a table of tensor with the range of parameters
123     to use
124     — Current implementation for 3 features only
125     — alphas_table = {alpha_w_tensor, alpha_c_tensor}
126     — Return a tensor with first columns the alpha value and last the
127     score for each
128     local nfeatures = #alphas_table
129     local v_seq_dev, precision, recall, f
130     local alphas = torch.DoubleTensor(1+nfeatures)
131     local size1 = alphas_table[1]:size(1)
132     local size2 = alphas_table[2]:size(1)
133     local size3 = alphas_table[3]:size(1)
134     local num_evaluations = size1*size2*size3
135     local score_ind = 1
136
137     — Columns for 2 features are (alphas_w_value, alphas_c_value,
138     f_score, precision, recall)
139     local scores = torch.DoubleTensor(num_evaluations, nfeatures+3)

```

```

137     for i=1,size1 do
138         alpha_w = alphas_table[1][i]
139         for k=1,size2 do
140             alpha_c = alphas_table[2][k]
141             for j=1,size3 do
142                 alpha_p = alphas_table[3][j]
143                 alphas:copy(torch.Tensor({alpha_t, alpha_w, alpha_c,
144                                         alpha_p}))
145                 v_seq_dev = predict(observations, emissions, transition
146                                     , alphas, nfeatures)
147                 precision, recall = compute_score(v_seq_dev,
148                                                     true_classes)
149                 f = f_score(precision, recall)
150                 — Filling the scores tensor
151                 scores[{score_ind, 1}] = alpha_w
152                 scores[{score_ind, 2}] = alpha_c
153                 scores[{score_ind, 3}] = alpha_p
154                 scores[{score_ind, 4}] = f
155                 scores[{score_ind, 5}] = precision
156                 scores[{score_ind, 6}] = recall
157                 score_ind = score_ind + 1
158             end
159         end
160     end
161
162     return scores
163 end
164
165 function main()
166     — Parse input params
167     opt = cmd:parse(arg)
168
169     — Reading file from pre-processing
170     myFile = hdf5.open(opt.datafile, 'r')
171     data = myFile:all()
172     emission_w = data['emission_w']
173     emission_c = data['emission_c']
174     emission_p = data['emission_p']
175     print(emission_p:size())
176     — Table of emission tensor (one tensor per feature)
177     emissions = {emission_w, emission_c, emission_p}
178     — Assertion on number of features
179     nfeatures = opt.nfeatures
180     if nfeatures > #emissions then

```

```

179         error('Too many features specified ')
180     end
181     print('Number of features used: '..nfeatures)
182     transition = data['transition ']
183     input_matrix_train = data['input_matrix_train ']
184     input_matrix_dev = data['input_matrix_dev ']
185     input_matrix_test = data['input_matrix_test ']
186     myFile:close()
187
188     — Parameters:
189     true_classes = input_matrix_dev:narrow(2,6,1):clone():view(
        input_matrix_dev:size(1))
190     — contain in each column feature observation
191     — (same order as the feature emission tensor in the emissoins
        table)
192     observations = input_matrix_dev:narrow(2,3,nfeatures):clone()
193     — Alpha parameter
194     alphas = torch.Tensor({opt.alpha_t, opt.alpha_w, opt.alpha_c, opt.
        alpha_p })
195
196     — Prediction on dev
197     v_seq_dev = predict(observations, emissions, transition, alphas,
        nfeatures)
198     print(v_seq_dev:size(1))
199     precision, recall = compute_score(v_seq_dev, true_classes)
200     f = f_score(precision, recall)
201
202     print('Prediction on dev')
203     print('Precision is : '..precision)
204     print('Recall is : '..recall)
205     print('F score (beta = 1) is : '..f)
206
207     — Cross validation
208     if (opt.cv == 1) then
209         alphas_table = {}
210         — alpha_w
211         alphas_table[1] = torch.Tensor({0.1, 0.2, 0.3, 0.4, 0.5})
212         — alpha_c
213         alphas_table[2] = torch.Tensor({5, 8, 10, 12})
214         — alpha_p
215         alphas_table[3] = torch.Tensor({1, 2, 4, 6})
216
217         scores = cross_validation(observations, emissions, transitions,
            true_classes,
218
            alphas_table, opt.alpha_t)

```

```

219         print(scores)
220
221     — Saving the score
222     myFile = hdf5.open('plot_scores.hdf5', 'w')
223     myFile:write('scores', scores)
224     myFile:close()
225     print('CV on dev saved in '..plot_scores.hdf5')
226 end
227
228 — Prediction on test
229 if (opt.test == 1) then
230     print('Prediction on test')
231     observations_test = input_matrix_test:narrow(2,3,nfeatures):
        clone()
232     v_seq_test = predict(observations_test, emissions, transition,
        alphas, nfeatures)
233     — Saving predicted sequence on test
234     myFile = hdf5.open(opt.datafile_test, 'w')
235     myFile:write('v_seq_test', v_seq_test)
236     myFile:write('v_seq_dev', v_seq_dev)
237     myFile:close()
238     print('Sequence predicted on test saved in '..opt.datafile_test
        )
239 end
240
241 end
242
243 main()

```

Max-Entropy Markov Model:

```

1  require 'hdf5';
2  require 'nn';
3  require 'helper.lua';
4
5  — Loading data
6  myFile = hdf5.open('../data/MM_data_pos.hdf5', 'r')
7  data = myFile:all()
8  input_matrix_train_pos = data['input_matrix_train_pos']
9  input_matrix_dev_pos = data['input_matrix_dev_pos']
10 input_matrix_test_pos = data['input_matrix_test_pos']
11 myFile:close()
12
13 nwords = input_matrix_train_pos:size(1)
14 train_output = input_matrix_train_pos:narrow(2,59)
15 train_input_pos = torch.Tensor(nwords-1,1+9+5+43)

```

```

16 train_input_pos:narrow(2,1,1):copy(input_matrix_train_pos:narrow(2,1,1)
    :narrow(1,2,nwords-1))
17 train_input_pos:narrow(2,2,9):copy(input_matrix_train_pos:narrow(2,2,9)
    :narrow(1,1,nwords-1))
18 train_input_pos:narrow(2,11,5):copy(input_matrix_train_pos:narrow
    (2,11,5):narrow(1,1,nwords-1))
19 train_input_pos:narrow(2,16,43):copy(input_matrix_train_pos:narrow
    (2,16,43):narrow(1,1,nwords-1))
20
21 observations_dev = input_matrix_dev_pos:narrow(2,1,1):clone()
22 dev_feat = input_matrix_dev_pos:narrow(2,11, 5 + 43)
23 dev_true_classes = input_matrix_dev_pos:narrow(2, 59,1):squeeze()
24
25 observations_test_pos = input_matrix_test_pos:narrow(2,1,1)
26 observations_test_feat = input_matrix_test_pos:narrow(2,2,5+43)
27
28 — Defining the model
29
30 model = nn.Sequential()
31 t1_pos = nn.ParallelTable()
32
33 t1_pos_1 = nn.Sequential()
34 t1_pos_1:add(LT)
35 t1_pos_1:add(nn.View(-1,50))
36
37 t1_pos_2 = nn.Identity()
38
39 t1_pos:add(t1_pos_1)
40 t1_pos:add(t1_pos_2)
41
42 model:add(t1_pos)
43 model:add(nn.JoinTable(2))
44
45 model:add(nn.Linear(50 + 9 + 5 + 43,9))
46 model:add(nn.LogSoftMax())
47
48 — Training function:
49
50
51 function train_model_cap(train_input, train_output, model, criterion,
    din, nclass, eta, nEpochs, batchSize)
52     — Train the model with a mini batch SGD
53     — standard parameters are
54     — nEpochs = 1
55     — batchSize = 32

```

```

56 — eta = 0.01
57 local loss = torch.Tensor(nEpochs)
58
59 — To store the loss
60 local av_L = 0
61
62 — Memory allocation
63 local inputs_batch = torch.DoubleTensor(batchSize, din)
64 local targets_batch = torch.DoubleTensor(batchSize)
65 local outputs = torch.DoubleTensor(batchSize, nclass)
66 local df_do = torch.DoubleTensor(batchSize, nclass)
67
68 for i = 1, nEpochs do
69     — timing the epoch
70     timer = torch.Timer()
71     av_L = 0
72
73     — mini batch loop
74     for t = 1, train_input:size(1), batchSize do
75         — Mini batch data
76         current_batch_size = math.min(batchSize, train_input:size(1)
77             -t)
78
79         inputs_batch:narrow(1,1,current_batch_size):copy(
80             train_input:narrow(1,t,current_batch_size))
81
82         targets_batch:narrow(1,1,current_batch_size):copy(
83             train_output:narrow(1,t,current_batch_size))
84
85         — reset gradients
86         model:zeroGradParameters()
87
88         — Forward pass (selection of inputs_batch in case the
89             batch is not full, ie last batch)
90         outputs:narrow(1,1,current_batch_size):copy(model:forward({
91             inputs_batch:narrow(1,1,current_batch_size):narrow
92                 (2,1,1),
93             inputs_batch:narrow(1,1,current_batch_size):narrow(2,2,din
94                 -1)}))
95
96         — Average loss computation
97         f = criterion:forward(outputs:narrow(1,1,current_batch_size
98             ), targets_batch:narrow(1,1,current_batch_size))
99
100         av_L = av_L +f

```

```

93         — Backward pass
94         df_do:narrow(1,1,current_batch_size):copy(criterion:
            backward(outputs:narrow(1,1,current_batch_size),
            targets_batch:narrow(1,1,current_batch_size)))
95         model:backward({inputs_batch:narrow(1,1,current_batch_size)
            :narrow(2,1,1), inputs_batch:narrow(1,1,
            current_batch_size):narrow(2,2,din-1)},
            df_do:narrow(1,1,current_batch_size))
96
97
98         model:updateParameters(eta)
99
100     end
101
102     print('Epoch '..i..'': '..timer:time().real)
103     loss[i] = av_L/math.floor(train_input:size(1)/batchSize)
104     print('Average Loss: '.. loss[i])
105
106 end
107
108 return loss
109 end
110
111 — Viterbi for MEMM:
112
113 — Evaluates the matrix of scores for all possible tags for the
    previous word, using the word features at timestep i
114
115 function compute_logscore_extrafeat(observations, feat, i, model, C)
116     local y = torch.zeros(C,C)
117     local hot_1 = torch.zeros(C+feat:size(2))
118     for j = 1, C do
119         hot_1:zero()
120         hot_1[j] = 1
121         hot_1:narrow(1,10,feat:size(2)):copy(feat:narrow(1,i,1))
122         y:narrow(1,j,1):copy(model:forward({observations[i]:view(1,1),
            hot_1:view(1,C+feat:size(2))}))
123     end
124     return y
125 end
126
127 — Evaluates the highest scoring sequence:
128 function viterbi_extrafeat(observations, feat, compute_logscore, model,
    C)
129
130     local y = torch.zeros(C,C)

```

```

131 — Formating tensors
132 local initial = torch.zeros(C, 1)
133 — initial started with a start of sentence: <t>
134
135 initial[{8,1}] = 1
136 initial:log()
137
138 — number of classes
139 local n = observations:size(1)
140 local max_table = torch.Tensor(n, C)
141 local backpointer_table = torch.Tensor(n, C)
142 — first timestep
143 — the initial most likely paths are the initial state distribution
144 local maxes, backpointers = (initial + compute_logscore_extrafeat(
    observations, feat, 1, model, C)[8]):max(2)
145 max_table[1] = maxes
146 — remaining timesteps ("forwarding" the maxes)
147 for i=2,n do
148     — precompute edge scores
149
150     y:copy(compute_logscore_extrafeat(observations, feat, i, model,
        C))
151     scores = y:transpose(1,2) + maxes:view(1, C):expand(C, C)
152
153     — compute new maxes
154     maxes, backpointers = scores:max(2)
155
156     — record
157     max_table[i] = maxes
158     backpointer_table[i] = backpointers
159 end
160 — follow backpointers to recover max path
161 local classes = torch.Tensor(n)
162 maxes, classes[n] = maxes:max(1)
163 for i=n,2,-1 do
164     classes[i-1] = backpointer_table[{i, classes[i]}]
165 end
166
167 return classes
168 end
169
170 — Train Model
171
172 loss_pos = train_model_cap(train_input_pos, train_output,
    ultimate_t_pos, criterion, 1 + 9 + 5 + 43, 9, 0.1, 20, 32)

```



```

173
174 — Evaluate performance on dev set:
175
176 cl_pos_dev = viterbi_extrafeat(observations_dev, dev_feat,
    compute_logscore_extrafeat, ultimate_t_pos, 9)
177 f = f_score(cl_pos_dev, dev_true_classes)
178
179 — Predict on test:
180
181 v_seq_test_pos = viterbi_extrafeat(observations_test_pos,
    observations_test_feat, compute_logscore_extrafeat, ultimate_t_pos,
    9)
182
183 — Saving predicted sequence on test
184 myFile = hdf5.open('../submission/v_seq_test_mem_pos', 'w')
185 myFile.write('v_seq_test', v_seq_test_pos)
186 myFile.close()

```

Structured Perceptron:

```

1 function train_model(train_input, sent, train_output, observations_dev,
    model, din, nclass, eta, nEpochs)
2     — Train the model with the structured perceptron approach
3     — V1: only treating the eged leaving the error
4
5     — For the verbose print
6     observations = observations_dev:narrow(2,1,1):narrow(1,1,1000):
        clone()
7     true_classes = observations_dev:narrow(2,16,1):narrow(1,1,1000):
        squeeze()
8
9     — Memory allocation
10    inputs_batch = torch.DoubleTensor(100, din)
11    gold_sequence = torch.DoubleTensor(100)
12    high_score_seq = torch.DoubleTensor(100)
13    grad_pos = torch.zeros(9)
14    grad_neg = torch.zeros(9)
15    pr1 = torch.zeros(9)
16    pr2 = torch.zeros(9)
17
18    for i = 1, nEpochs do
19        — timing the epoch
20        timer = torch.Timer()
21
22        — mini batch loop
23        for t = 2, sent:size(1)-1 do

```

```

24      — Mini batch data
25      sent_size = sent[{t,2}]
26  —      print('here1 ')
27
28      inputs_batch:narrow(1,1,sent_size+1):copy(train_input:
29      narrow(1,sent[{t,1}]-1,sent_size+1))
30      print('here2 ')
31
32      gold_sequence:narrow(1,1,sent_size+1):copy(train_output:
33      narrow(1,sent[{t,1}]-1,sent_size+1))
34      print('here3 ')
35
36      — reset gradients
37      model:zeroGradParameters()
38      —gradParameters:zero()
39
40      — Forward pass on a batch subsequence:
41      high_score_seq:narrow(1,1,sent_size+1):copy(viterbi(
42      inputs_batch:narrow(1,1,sent_size+1):narrow(2,1,1),
43      compute_logscore
44      ,
45      model
46      ,
47      nclass
48      ))
49
50      print('here4 ')
51
52      for ii = 1, sent_size+1 do
53      grad_pos:zero()
54      if high_score_seq[ii] ~= gold_sequence[ii] then
55      — WARNING: Need to call backward right after the
56      forward with the same input to compute correct
57      gradients
58
59      — Use of a single gradient (grad_pos) with a
60      penalization on the wrong class predicted (1)
61      — and a valorisation (-1) on the correct class to
62      predict
63      — We treat here only the transition after the
64      error
65      model:forward({inputs_batch:narrow(1,ii,1):narrow
66      (2,1,1),inputs_batch:narrow(1,ii,1):narrow
67      (2,2,9)})
68      grad_pos[gold_sequence[ii]] = -1

```

```

54         grad_pos[high_score_seq[ii]] = 1
55         model:backward({ inputs_batch:narrow(1,ii,1):narrow
                        (2,1,1), inputs_batch:narrow(1,ii,1):narrow
                        (2,2,9)}, grad_pos:view(1,9))
56
57
58         end
59     end
60 —     print('here7')
61     model:updateParameters(eta)
62
63     end
64
65     print('Epoch '..i..'': '..timer:time().real)
66     — Print the f-score on a the first 1000 words to follow the
        improvement of the model
67     cl = viterbi(observations, compute_logscore, model, 9)
68     print (f_score(cl, true_classes))
69
70     end
71 end
72
73 function train_model2(train_input, sent, train_output, model, din,
        nclass, eta, nEpochs, obs_val, true_val, f_score)
74     — Train the model with the structured perceptron approach
75     — V2: treating the two edges, the one leading to the error and the
76     — one leaving the error.
77
78     val_res = torch.zeros(nEpochs,3)
79     — Memory allocation
80     inputs_batch = torch.DoubleTensor(100, din)
81     gold_sequence = torch.DoubleTensor(100)
82     high_score_seq = torch.DoubleTensor(100)
83     grad_pos = torch.zeros(9)
84     grad_neg = torch.zeros(9)
85     one_hot_true = torch.zeros(1,9)
86     one_hot_false = torch.zeros(1,9)
87
88     for i = 1, nEpochs do
89         — timing the epoch
90         timer = torch.Timer()
91
92         — mini batch loop
93         for t = 2, sent:size(1)-1 do
94             — Mini batch data

```

```

95         sent_size = sent[{t,2}]
96 —         print('here1 ')
97
98         inputs_batch:narrow(1,1,sent_size+1):copy(train_input:
          narrow(1,sent[{t,1}]-1,sent_size+1))
99 —         print('here2 ')
100
101         gold_sequence:narrow(1,1,sent_size+1):copy(train_output:
          narrow(1,sent[{t,1}]-1,sent_size+1))
102 —         print('here3 ')
103
104         — reset gradients
105         model:zeroGradParameters()
106         —gradParameters:zero()
107
108         — Forward pass on a batch subsequence:
109         high_score_seq:narrow(1,1,sent_size+1):copy(viterbi(
          inputs_batch:narrow(1,1,sent_size+1):narrow(2,1,1),
110                                                     compute_logscore
                                                         '
                                                         model
                                                         '
                                                         nclass
                                                         ))
111 —         print('here4 ')
112
113         previous_error = false
114
115         for ii = 1, sent_size+1 do
116
117             grad_neg:zero()
118             grad_pos:zero()
119
120             if high_score_seq[ii] ~= gold_sequence[ii] and not
              previous_error then
121                 — WARNING: Need to call backward right after the
                  forward with the same input to compute correct
                  gradients
122
123                 — Use of a single gradient (grad_pos) with a
                  penalization on the wrong class predicted (1)
124                 — and a valorisation (-1) on the correct class to
                  predict
125
126                 model:forward({inputs_batch:narrow(1,ii,1):narrow

```

```

(2,1,1), inputs_batch:narrow(1, ii, 1):narrow
(2,2,9))
127 grad_pos[ gold_sequence[ ii ]] = -1
128 grad_pos[ high_score_seq[ ii ]] = 1
129 model:backward({ inputs_batch:narrow(1, ii, 1):narrow
(2,1,1), inputs_batch:narrow(1, ii, 1):narrow
(2,2,9) }, grad_pos:view(1,9))

130
131 grad_neg:zero()
132 grad_pos:zero()
133 if ii ~= (sent_size + 1) then
134     one_hot_true:zero()
135     one_hot_true[1][ gold_sequence[ ii ]] = 1
136     model:forward({ inputs_batch:narrow(1, ii+1,1):
narrow(2,1,1), one_hot_true })
137     grad_neg[ gold_sequence[ ii+1 ]] = -1
138     model:backward({ inputs_batch:narrow(1, ii+1,1):
narrow(2,1,1), one_hot_true }, grad_neg:view
(1,9) )

139
140     one_hot_false:zero()
141     one_hot_false[1][ high_score_seq[ ii ]] = 1
142     model:forward({ inputs_batch:narrow(1, ii+1,1):
narrow(2,1,1), one_hot_false })
143     grad_pos[ gold_sequence[ ii+1 ]] = 1
144     model:backward({ inputs_batch:narrow(1, ii+1,1):
narrow(2,1,1), one_hot_false }, grad_pos:view
(1,9) )

145 end
146
147 previous_error = true
148
149 elseif high_score_seq[ ii ] ~= gold_sequence[ ii ] and
previous_error then
150
151     if ii ~= sent_size + 1 then
152         one_hot_true:zero()
153         one_hot_true[1][ gold_sequence[ ii ]] = 1
154         model:forward({ inputs_batch:narrow(1, ii+1,1):
narrow(2,1,1), one_hot_true })
155         grad_neg[ gold_sequence[ ii+1 ]] = -1
156         model:backward({ inputs_batch:narrow(1, ii+1,1):
narrow(2,1,1), one_hot_true }, grad_neg:view
(1,9) )
157

```

```

158             one_hot_false:zero()
159             one_hot_false[1][high_score_seq[ii]] = 1
160             model:forward({inputs_batch:narrow(1,ii+1,1):
                           narrow(2,1,1),one_hot_false})
161             grad_pos[gold_sequence[ii+1]] = 1
162             model:backward({inputs_batch:narrow(1,ii+1,1):
                           narrow(2,1,1),one_hot_false}, grad_pos:view
                           (1,9) )
163         end
164
165         previous_error = true
166
167         else
168             previous_error = false
169         end
170     end
171 —     print('here7')
172     model:updateParameters(eta)
173
174 end
175
176 print('Epoch '..i..'': '..timer:time().real)
177 cl = viterbi(obs_val, compute_logscore, model, 9)
178 val_res[i][1], val_res[i][2], val_res[i][3] = f_score(cl,
               true_val)
179 print('f-score: '.. val_res[i][1])
180
181 end
182 return val_res
183 end

```

Helper:

```

1 — function to evaluate the predicted sequence
2 — need to compute precision and recall (class 1 stands for negative
   class)
3 function compute_score(predicted_classes, true_classes)
4     print('here')
5     local n = predicted_classes:size(1)
6     local right_pred = 0
7     local positive_true = 0
8     local positive_pred = 0
9     for i=1,n do
10         if predicted_classes[i] > 1 then
11             positive_pred = positive_pred + 1
12         end

```

```
13         if true_classes[i] > 1 then
14             positive_true = positive_true + 1
15         end
16         if (true_classes[i] == predicted_classes[i]) and true_classes[i]
17             ] > 1 then
18             right_pred = right_pred + 1
19         end
20     end
21     local precision = right_pred/positive_pred
22     local recall = right_pred/positive_true
23     return precision , recall
24 end
25 function f_score(precision , recall)
26     return 2*precision*recall/(precision+recall)
27 end
```